Model Risk Remediation Plan

Document ID: MAP-RP-001-COMPANION
Version: 1.0
Status: Final — Portfolio Artifact
Author: Milton Wani
Date: November 2026

---

1. Purpose of This Document

This document accompanies the Model Risk Remediation Plan (MAP-RP-001) [GitHub]. It explains, in plain language, what the remediation plan contains, why it was built the way it was, and what it demonstrates about model risk governance practice.

The remediation plan itself is a structured artifact: a findings register, a root cause analysis, a tiering and prioritization framework, a 6-month roadmap, compensating controls, and a closure and sign-off process. This companion document tells the story behind that structure.

The intended reader is a hiring manager, a model risk practitioner, or a reviewer who wants to understand the thinking behind the artifact, not just the artifact itself.

---

2. What Problem This Plan Solves

When a model validation raises findings, the validation report is only the beginning of the work. The harder and more consequential work is what happens next: who owns the fix, how the fix will be evidenced, when it must be closed, and how an independent reviewer confirms that closure is real rather than self-attested.

This plan addresses that gap. It is not a validation opinion. It does not authorize or reject any model. It is the **audit trail between a finding being raised and a finding being closed**.

Ten findings are in scope, selected from two validation pipelines (Lending Club credit model and Fannie Mae IFRS 9 stress testing). The selection criteria were:

- **Severity:** CRITICAL or HIGH in the source registers.
- **Materiality:** the finding affects the deployability of the model it touches.
- **Closeability:** the finding can be closed within a 6-month horizon with identifiable evidence.

Findings that are documentation-only, or that require data the project does not have, are recorded in the validation opinions but excluded from this plan. They belong in a longer-horizon roadmap.

3. How the Plan Is Structured

The plan follows a logical sequence that mirrors how remediation is governed in a regulated institution.

3.1 Findings at a Glance

A single table gives the reviewer the full scope in one view: finding ID, source project, severity, tier, SLA, and status. The table is explicitly framed as a **reading aid, not a status tracker**. The live tracker is the MRM dashboard described in Section 7 of the plan.

3.2 Issue Identification and Root Cause Analysis

For each finding, the plan records:

- The **root cause category** (conceptual, methodology, calibration, discrimination, data/scope, scenario calibration, parameter stability).
- A **root cause detail** explaining what actually went wrong.

The root cause analysis is the most important section of the plan. It distinguishes between:

- **Coding defects** (fixable by the model development team).
- **Structural limitations** (data the project does not have; scope the project cannot cover).
- **Methodology errors** (the model was built differently from its documented design).

This distinction matters because it determines whether a finding is closeable within the horizon or whether it must be documented as a permanent limitation.

3.3 Model Tiering and Prioritization

Each finding is assigned a tier (1–3) based on:

- **Financial materiality** (what happens if the model is wrong).
- **Complexity** (how hard the fix is).
- **Usage scale** (how many decisions the model affects).

The tier assignment reflects the tier the model **would carry if promoted to its stated intended use**, not its current demonstration status. This is a deliberate choice: it forces the plan to treat the findings with the seriousness they would receive in production.

The prioritization order for the 6-month roadmap follows from the tiering. CRITICAL and Tier 1 findings are addressed first. Structural limitations that require data acquisition are placed later in the horizon.

3.4 Remediation Actions and Milestones

For each finding, the plan specifies:

- The **action** (what will be done).
- The **owner** (who is accountable).
- The **milestone** (when it will be done, relative to plan approval).
- The **closure evidence** (what artifact proves the fix is real).

The closure evidence column is the most important. A finding cannot be closed by assertion. It is closed by a new pipeline run, a new validation report, or a signed limitation. This is the discipline that separates a remediation plan from a to-do list.

3.5 Compensating Controls

While permanent fixes are under development, compensating controls limit the damage the finding can cause. Examples:

- If a model has a CRITICAL fairness finding, it may not be used for any consequential decision. Scores may be used for research only.
- If a split methodology understates out-of-time degradation, all metrics are labeled as in-sample optimistic.
- If an EAD ratio is outside band, ECL numbers are not cited until the ratio is resolved.

Compensating controls are not a substitute for remediation. They are a bridge that allows the organization to continue operating safely while the fix is built.

3.6 Independent Re-Validation and Sign-Off

Each finding's closure requires independent re-validation, not self-attestation by the remediation owner. The re-validation must be performed by a reviewer who did not implement the fix, and the evidence must be a new pipeline run.

The closure evidence template requires:

1. Finding ID and severity.
2. Root cause (from Section 2).
3. Action taken (from Section 4).
4. New run ID and artifact hash.
5. Before/after comparison.
6. Residual risk.
7. Independent reviewer name, role, date, signature.
8. Closure approval.

The plan deliberately leaves the sign-off fields as [ANALYST: insert] placeholders. This is not an omission. It is a statement that **the plan produces the structure; the analyst produces the closure evidence and the signature**. A remediation plan that auto-fills its own sign-off is not an independent re-validation.

3.7 Executive Tracking and Reporting

The plan defines the fields the MRM dashboard must track per finding:

- Status (OPEN / IN PROGRESS / CLOSED / DEFERRED).
- Days open.
- Days remaining to SLA.
- SLA breach flag.
- Escalation flag.
- Owner.
- Evidence artifact link.

Monthly reporting to the Model Risk Committee covers findings closed, findings slipping past SLA, findings with residual risk above threshold, and any finding where the remediation action has changed scope. Board-level visibility is reserved for CRITICAL findings and for any finding where SLA is breached by more than 30 days.

---

4. What This Plan Demonstrates

This artifact demonstrates several capabilities that are directly relevant to model risk governance and GRC roles.

Capability	Where It Appears
Root cause analysis	Section 2 of the plan: distinguishing coding defects from structural limitations and methodology errors.
Risk tiering and prioritization	Section 3 of the plan: assigning tiers based on materiality, complexity, and usage scale.
Remediation design	Section 4 of the plan: action, owner, milestone, and closure evidence for each finding.
Compensating controls	Section 5 of the plan: bridge controls while permanent fixes are built.
Independent re-validation	Section 6 of the plan: closure requires a new run, not a narrative claim.
Executive reporting	Section 7 of the plan: dashboard fields and monthly MRC reporting.
Regulatory alignment	The plan references OSFI E-23, SR 11-7, IFRS 9, and ECOA as the frameworks the findings are mapped against.

5. How to Read the Plan

If you are a hiring manager or a reviewer, here is the fastest path through the document.

1. **Start with the findings-at-a-glance table.** It gives you the scope in one view.
2. **Read the root cause analysis for the CRITICAL finding (LC-R-01).** It shows how the plan distinguishes a proxy-validity problem from a coding defect.
3. **Read the tiering rationale for the two Tier 1 findings (LC-R-01 and LC-R-05).** It shows how the plan thinks about materiality and structural limitation.
4. **Read the closure evidence column for any finding.** It shows what proof of closure looks like.
5. **Read the sign-off note.** It shows that the plan is designed for independent re-validation, not self-attestation.

6. Scope and Limitations

This plan is a portfolio artifact. It is built on two independent study projects (Lending Club credit model validation and Fannie Mae IFRS 9 stress testing), not on a production model at a regulated institution.

The findings are real findings within those projects. The remediation actions are designed to be closeable within a 6-month horizon. The closure evidence requirements are designed to be independently verifiable.

The plan does not constitute legal advice, regulatory interpretation, or a validation opinion. It is a demonstration of remediation governance methodology.